

A Million Ways to Fail in Production

Embracing Catastrophes for Fun and Profit

Jonatan Ivanov
2025-02-28



About Me

- Spring Team
 - Micrometer
 - Spring Cloud, Spring Boot
 - Spring Observability Team
- Java Champion
- Seattle Java User Group
- develotters.com
- **@jonatan_ivanov**



What can go wrong? 🙄

What can go wrong?

- You get an error (something that contains an error)
- Or nothing (timeout, connection reset, TLS, etc.)
- Latency
- Fallacies of distributed computing
- Network, Hardware, Apps, DBs, Caches, Streams, Queues, etc.

Why do we care?

- We need to fix them
- Increasing complexity
- Something is always broken on huge scale
- Chaotic environments
- Unknown unknowns
- Things can be perceived differently by observers

What can we do about it?

Observability: Data about your components

- Logging: What happened? (Why?)
- Metrics: What's the context? How bad is it?
- Distributed Tracing: Why did it happen?
- But there is so much more...

Memory leak only in PROD

Memory leak only in PROD

- The platform/container was homegrown 🤖
- Had memory leak on Java 5 💣
- It was fine on Java 6 😎 💣
- We deployed new features (automated) 🚚
- Also updated Java: 6u123 → 6u124 ☕
- The app was leaking! 😱
- But only in prod! 🕵️
- Investigations for weeks! 🤔

What happened?

- I “accidentally” ran **java -version**
- It said Java 5 ...
- The 6u124 folder contained Java 5
- The deployment script used that
- The app was running on Java 5
- And leaking ...

What can we learn from it?








- Do not write platforms/containers on your own
- Do not 100% trust your deployment pipeline
- Only your app knows its environment
 - Java version, vendor, ...
 - OS version, arch, ...
 - ENV vars, properties, ...

Solution: Ask the app!

```
"java": {  
  "version": "23",  
  "vendor": {  
    "name": "BellSoft"  
  },  
  "jvm": {  
    "name": "OpenJDK 64-Bit Server VM",  
    "vendor": "BellSoft",  
    "version": "23+38"  
  }  
}
```

Memory leak or not?

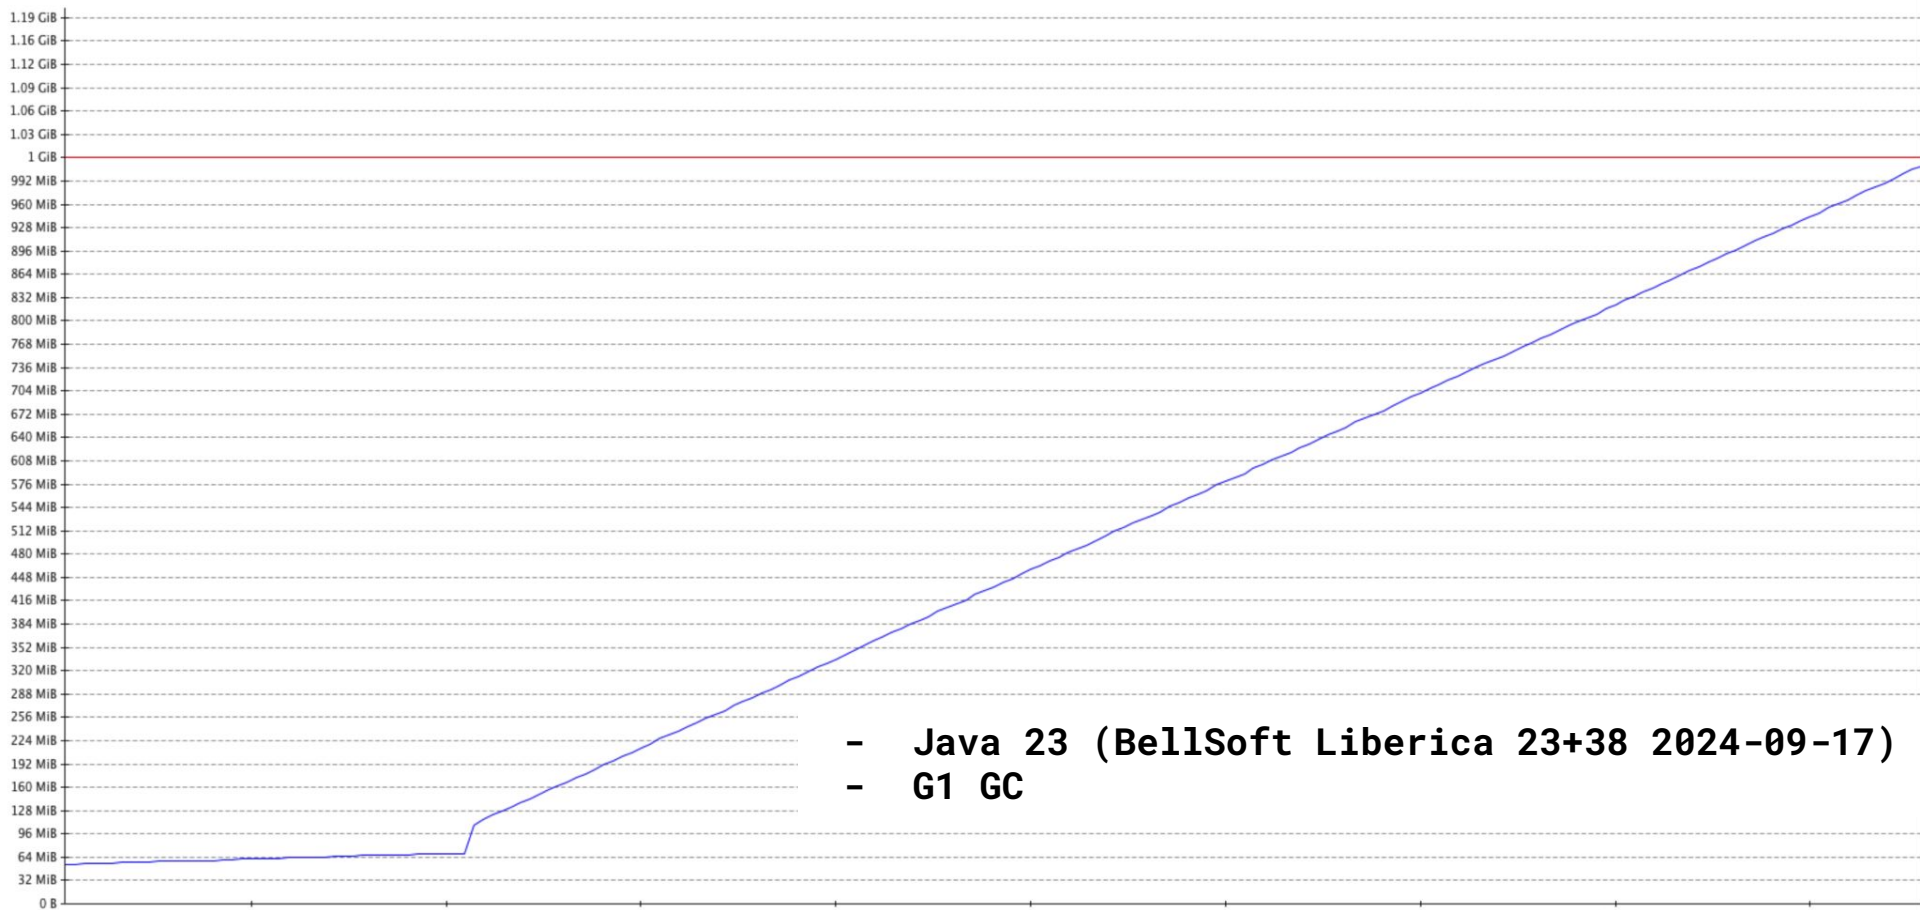
Memory leak or not?

- We deployed new features (manual steps) 
- One instance did not receive any traffic 
- No one noticed it 
- The app had a scheduled job (infrequent) 
- An alert was triggered! 
- The heap utilization was high (95+%)! 
- Only on the “no-traffic” instance! 
- Investigations, load tests...


What happened?

- A Generational GC was used...
- Let's play "Memory leak or not"...

Memory leak or not?



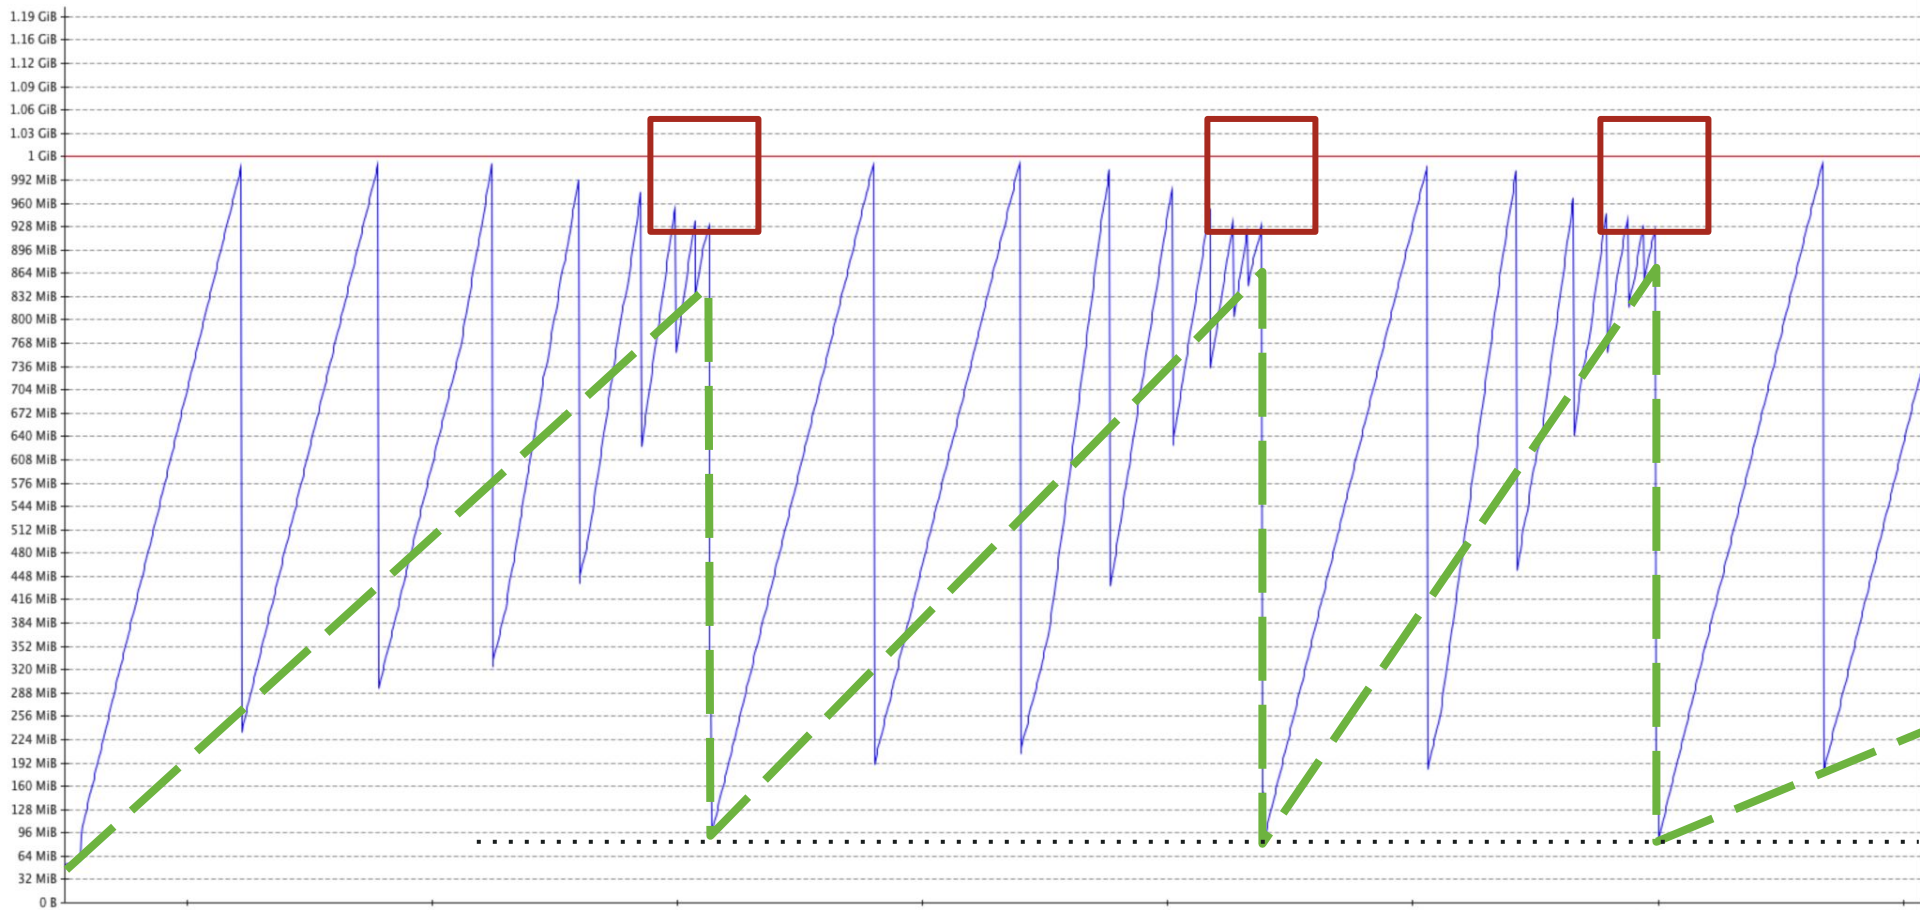
Memory leak or not?




Memory leak or not?



Memory leak or not?



Memory leak or not?



What happened?

- A Generational GC was used
- No “double-sawtooth pattern” in heap metrics
- No full GC (can’t say it’s a leak)
- Increase in heap usage was low (no traffic)
- So the GC was ok with high heap usage
- So didn’t trigger full GC (it’s expensive)
- Eventually it did
- Everything was back to normal

What can we learn from it?

- Human errors are inevitable
- GC's are complicated
- Even more complicated than that
- Nope, they are even more complicated

Solution

- Fully automate your deployments
- Observe and alert on traffic patterns
- Learn the JVM basics
- Observe heap usage
- Observe GC events

The app from the past

The app from the past

- We had a Service Registry (Eureka) 🎉
- An app was transferred to another team ➡️
- They removed the Eureka Client from it 💣
- Two years later the app appeared in Eureka 🕵️
- The team did not put the Eureka Client back 😱


What happened?

- A 2+ years old version of the app started
- The app was a data processing bot
- It processed data

What can we learn from it?

- We need more data about what is running!
- What apps are running (by env)?
- What versions are running (by env)?
- Where are they (by env)?
(host+port, instance, region, cloud provider, ...)
- How many instances (by env)?
- Service starts/stops (deployments, restarts)?

Solution: Service Registry

HOME LAST 1000 SINCE STARTUP

System Status

Environment	test	Current time	2024-09-11T23:09:13 +0000
Data center	default	Uptime	00:05
		Lease expiration enabled	true
		Renews threshold	3
		Renews (last min)	20

DS Replicas

[localhost](#)

Instances currently registered with Eureka

Application	AMIs	Availability Zones	Status
EUREKA	n/a (1)	(1)	UP (1) - 192.168.0.111:eureka:8761
SPRING-BOOT-ADMIN	n/a (1)	(1)	UP (1) - 192.168.0.111:spring-boot-admin:8080
TEA-SERVICE	n/a (1)	(1)	UP (1) - 192.168.0.111:tea-service:8090
TEALEAF-SERVICE	n/a (1)	(1)	UP (1) - 192.168.0.111:tealeaf-service:8092
WATER-SERVICE	n/a (1)	(1)	UP (1) - 192.168.0.111:water-service:8091

General Info

Name	Value
total-avail-memory	256mb
num-of-cpus	16

Solution: Service Registry

Spring Boot Admin

Wallboard

Applications

Journal

English

Applications


5

Instances

5






▼

Filter



all up

9/11/2024, 4:07:52 PM

>	<div>UP</div> EUREKA 1 instance	3cbc15c.1726095799979	
>	<div>UP</div> SPRING-BOOT-ADMIN 1 instance	3cbc15c.1726095799979	
>	<div>UP</div> TEA-SERVICE 1 instance	3cbc15c.1726095799979	
>	<div>UP</div> TEALEAF-SERVICE 1 instance	3cbc15c.1726095799979	
>	<div>UP</div> WATER-SERVICE 1 instance	3cbc15c.1726095799979	

The one you won't believe

The one you won't believe

- Trying to reproduce an intermittent issue (fixed)
- Running the same app in a loop 100 times
- Thousands of executions (start-work-verify-stop)
- One time, the app crashed 🤯

```
java.lang.ClassFormatError:  
Unknown constant tag 41 in class file  
io/rsocket/frame/FrameLengthCodec
```

What happened?

- **ClassFormatError**: the class file is malformed
- The class was loaded many times except once
- No dynamic class loading or byte-code generation
- Nothing was changed between executions
- No disk or memory issue was found
- Single-Event Upset (SEU)?

What can we learn from it?

- Anything that can go wrong, will go wrong
- Even if you think it cannot
- Unknown Unknowns
- Umwelt

Solution

- Observability

Expired TLS certificate

Expired TLS certificate

- What happened?
 - The certificate expired. 😓
- What can we learn?
 - Don't let the certificate expire! 😓
- Solution
 - Alert before it happens! ⚠️

(server/client certs, LB, API GW, etc.)

Ask your apps (and your LB, API GW, etc.)

```
"certificates": [  
  {  
    "subject": "CN=localhost,OU=Spring,L=Seattle,ST=WA,C=US",  
    "issuer": "CN=root,OU=Spring,L=Seattle,ST=WA,C=US",  
    "version": "V3",  
    "serialNumber": "64d019d1dd94eee0",  
    "signatureAlgorithmName": "SHA256withRSA",  
    "validityStarts": "2024-06-21T21:32:22Z",  
    "validityEnds": "2024-06-22T21:32:22Z",  
    "validity": {  
      "status": "WILL_EXPIRE_SOON",  
      "message": "Will expire within threshold (72h) at ..."  
    }  
  }  
]
```

What else?

- Kubernetes CPU rq and Memory limit
- Clock skew (or wrong timezone)
- Deploying the wrong version
- Deploying the wrong profile/config
- Unpatched OS
- Unpatched dependency (SBOM)
- Restarted the wrong instance

Thank you!

@jonatan_ivanov

develotters.com

slack.micrometer.io

GH: jonatan-ivanov/resourceater

