

Events, queues, projections, offline mode and much more



Wim Godden
Cu.be Solutions

Who am I ?

- Wim Godden

Where I'm from



Where I'm from



Where I'm from



Where I'm from



Where I'm from



Where I'm from



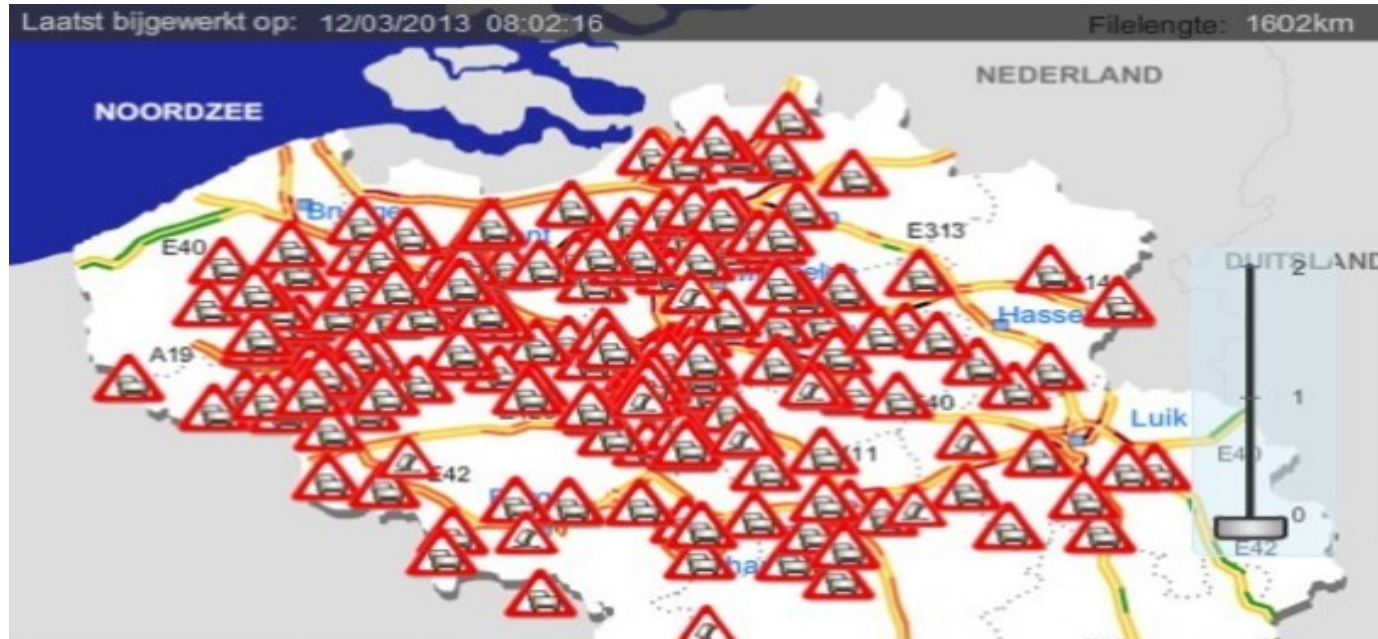
My town



My town



Belgium – the traffic



Who am I ?

- Wim Godden (@wimgtr)
- Founder of Cu.be Solutions (<https://cu.be>)
- Open Source developer since 1997
- Developer of PHPCompatibility, ...
- Speaker at PHP and Open Source conferences

Who are you ?

- Developers ?
- Devops ?
- Experience with event sourcing ?

Events, queues, projections and more

- Architecture
- How do you structure a project to grow in
 - Size
 - Complexity
 - Devices
- Understand what's going on
- Monitor

Short disclaimer

- Many possible approaches
- Let's look at one...
- ... which evolved over the years

One of our customers



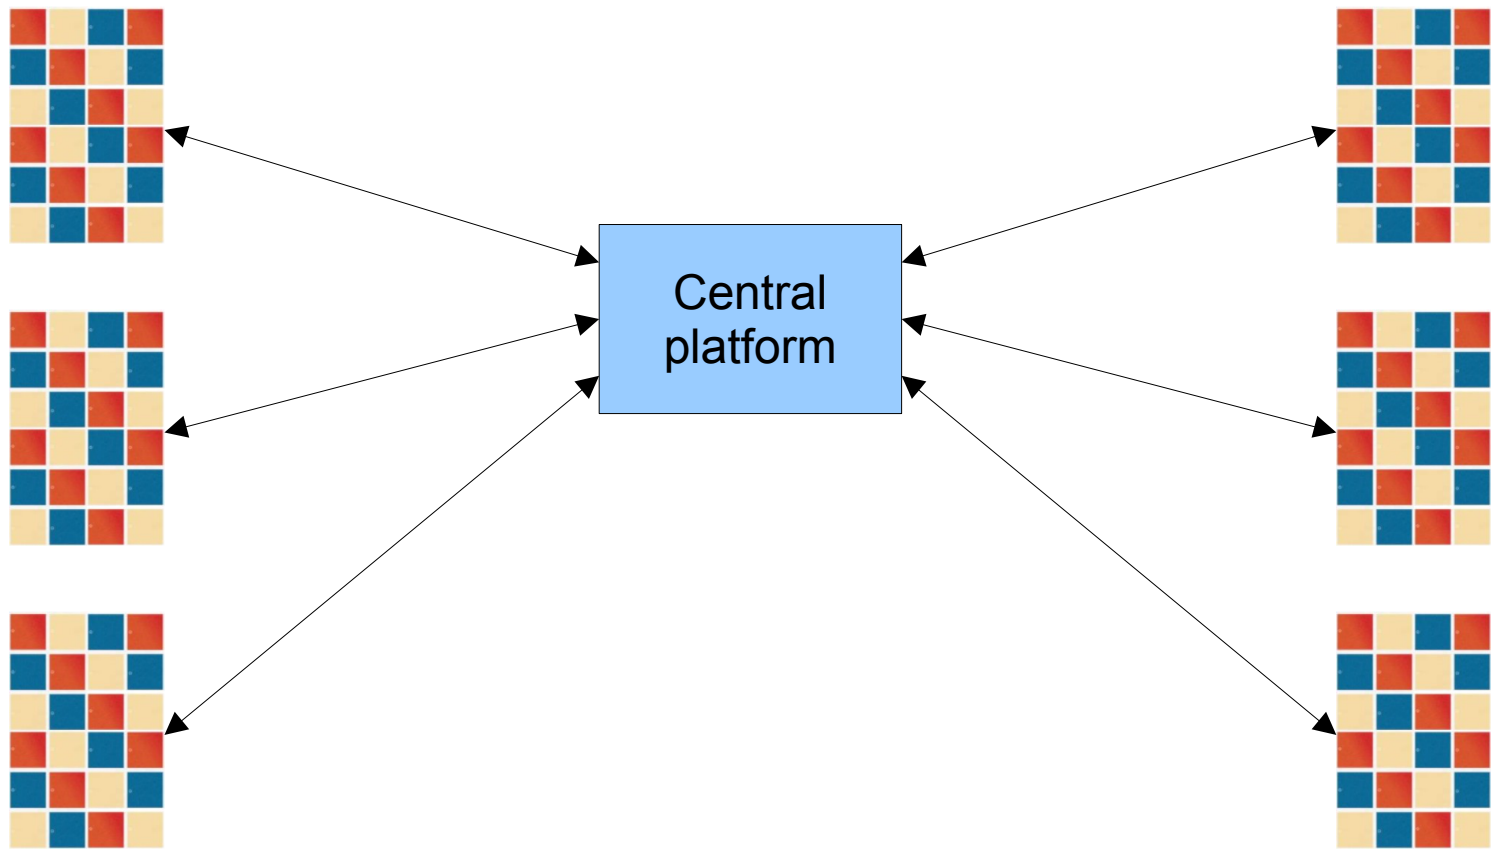
Many installations



These devices are in the field...



Central platform unites them all




Communication with all devices

- 1) Push configuration (prices, rental periods, ...)
- 2) Reserve lockers (through mobile app)
- 3) Receive events
- 4) ...

1. Pushing configuration to a unit

- Live or queued
- Live : API call → each unit has a unique name
- Queued : background worker attempts every
 - 60 sec
 - 5 min
 - 15 min
 - but only if the unit is online



How to detect it's online ?

<input type="checkbox"/>	NAME ↓	DISPLAY NAME	DEFAULT LOCKER UNITS	LAST RESPONSE TIME	LAST CONFIG PUSHED	LAST ONLINE AT
<input type="checkbox"/>	PL-1000	PL-1000	PL-1000	Offline	7 months ago	4 months ago
<input type="checkbox"/>	PL-1001	PL-1001	PL-1001	55 ms	a year ago	a minute ago
<input type="checkbox"/>	PL-1002	PL-1002	PL-1002	270 ms	4 months ago	a few seconds ago
<input type="checkbox"/>	PL-1003	PL-1003	PL-1003	275 ms	4 months ago	a few seconds ago
<input type="checkbox"/>	PL-1004	PL-1004	PL-1004	16 ms	3 months ago	a few seconds ago
<input type="checkbox"/>	PL-1005	PL-1005	PL-1005	16 ms	3 months ago	a few seconds ago
<input type="checkbox"/>	PL-1006	PL-1006	PL-1006	78 ms	3 months ago	a few seconds ago
<input type="checkbox"/>	PL-1007	PL-1007	PL-1007	Offline	never	10 days ago
<input type="checkbox"/>	PL-1008	PL-1008	PL-1008	Offline	never	10 days ago
<input type="checkbox"/>	PL-1009	PL-1009	PL-1009	39 ms	6 days ago	a few seconds ago

How to detect it's online ?

- 2 ways :
 - Let unit poll the central platform
 - Pro : Fast, doesn't put load on central platform
 - Con : Unit has to "inform" central platform that it's there
 - Let the central platform poll the unit
 - Pro : Central platform "knows" as soon as unit responds
 - Con :
 - Puts more load on central platform
 - Needs to scale...

Scaling process with little processing needs (like ping)

- No need to spin up multiple processes and a queue
- Instead : make the process asynchronous
 - use multiple threads
 - send pings to multiple devices at same time
- Unless...

Scaling to...

- 2017 : “Up to 100 over time”
- 2020 : Easily surpassed that...
- 2024 : “Should be ready for thousands”

How to distribute ?

- Geographically ?
 - Per continent ?
 - What about busy regions ?
- By device starting letter ?
 - Never a correct distribution
- Random ?
 - Requires an additional lookup table
- Solution : hash the device name, then distribute :
 - Initially by first letter
 - If you need more granularity later : by first 2 letters
 - etc.



2. Reserve / Rent a locker through mobile app

- Multiple API calls
 - Availability
 - Reserve
 - Confirm
 - Cancel
- Results in events being generated



Event sourcing


- Software architecture pattern
- Changes are stored as events
- Events define state

Event sourced system

- Single source of truth
- Sequence of events
- Stored in an event store
 - Could be table in DB
 - Apache Kafka
 - Redis

State

- current state = all events leading up to this moment



- future state = all events from this moment forward
 - Example : reservation of a future booking

Important rules

- Events are immutable
 - Append-only
 - No modifications
 - No deletions
 - Why ? Auditing
 - Theft
 - Accounting

Replaying events = state reconstruction

- Current (and future) state can be rebuilt
 - Just replay all events in order
- Bug in how state was determined ?
 - Fix the bug, replay all events

Note : If you have 100M+ events, this gets heavy ;-)

3. Receive events on central platform


Lots of events :

- Locker open/close
- Payment start, success, fail
- Hardware-related events (online, offline, error, ...)

How do we get them from the units to the central platform ?


Option A : fetch through API

- Poll every 60 seconds (if online)
- Ask for all new events
 - Update the last event id
- Store the events in the event store




Option A : fetch through API

- Poll every 60 seconds (if online)
- Ask for all new events
 - Update the last event id
- Store the events in the event store



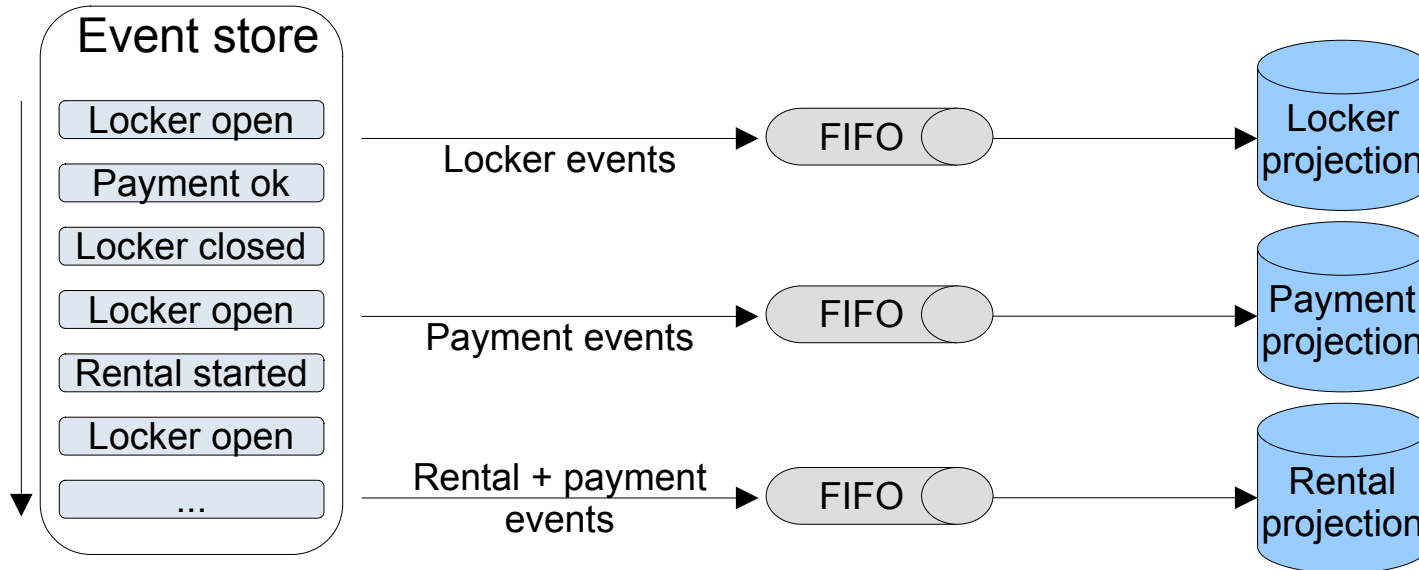
Option B : push to the central platform

- Push new events (grouped per x seconds)
- User experience = much better
- Careful though :
 - Load on central platform = dependent on external factors
 - High peaks can cause issues
 - Possible solution : use multiple projections to spread load



Creating projections

- Separate dataset
- Based on events
- Containing only the data needed for certain retrievals
- Goal : improve performance



Only need a small part of the data

Lockers / 8

Created At: 23 Apr 2024, 14:20 | Updated At: 12 Feb 2025, 10:51

<div>LockerEvent log</div>				
<div><div>Export Grid</div><div>Page: 8 of 15 Total of 354 records</div></div>				
LOCATION	OCCURRENCE	UNIT	OCCURRED	OCCURRED TIME ↑
Camping Festitent	NodesController/LockerClosed	A	7 months ago	21:34:59
Camping Festitent	NodesController/LockerOpen	A	7 months ago	21:34:46
Camping Festitent	NodesController/LockerClosed	A	7 months ago	12:17:44
Camping Festitent	NodesController/LockerOpen	A	7 months ago	12:16:54

Projections - pros/cons

- Pros :
 - Faster queries (compare to reading entire event log)
 - Perfect for summary dashboards
 - Bug in projection ? Fix it, then rebuild just 1 projection based on event log
 - Scalable and distributed
- Cons :
 - Asynchronous / eventual consistency → beware of race conditions
 - Requires extra CPU power
 - Data duplication

Online / offline mode

- Design for resilience
- In IOT / Mobile setup : design for poor / no connection
- Put your single source of truth locally
- Users can keep ordering via screen
- Events are stored locally
- Connection returns → events are pushed to central platform

Build it yourself / use existing tools ?

- Build it yourself :
 - Complete control
 - Not dependent on others
- Existing tools :
 - Leverage experience of others
 - Constant evolution and improvements
 - Plenty to choose from : Kafka, Pulsar, Redis Streams, ...

Questions ?



Questions ?

We're hiring @ Cu.be

Full Stack Engineers (Medior/Senior level)

→ <https://cu.be/jobs>

Contact

- Mastodon [@wimg@php.social](https://wimg@php.social)
- Slides <http://www.slideshare.net/wimg>
- E-mail wim@cu.be

Thanks !